

Critical Research Investigation: FreeRTOS Context Switch Failure on seL4 Microkernel

PhD Research - Secure Virtualization Systems

August 13, 2025

Abstract

This research note documents a critical blocking issue discovered during the integration of FreeRTOS with the seL4 formally verified microkernel. The investigation reveals a fundamental incompatibility between FreeRTOS's ARM context switching mechanism and seL4's virtualization environment, specifically related to the RFEIA (Return From Exception In ARM state) instruction. This issue prevents successful task execution and represents a significant barrier to running real-time operating systems under seL4 virtualization.

1 Problem Statement

1.1 Research Context

The integration of FreeRTOS (a commercial real-time operating system) with seL4 (a formally verified microkernel) represents a critical step toward achieving formally verified secure virtualization for real-time systems. However, our investigation has uncovered a fundamental compatibility issue that blocks successful execution.

1.2 Core Issue

FreeRTOS tasks fail to execute properly when running as a guest operating system under seL4's ARM virtualization framework. The system successfully completes all initialization phases but encounters a critical page fault during the first task context switch, specifically:

- **Fault Location:** PC: 0x8 (ARM Software Interrupt vector)
- **Fault Type:** Read prefetch fault
- **Context:** First task startup via `vPortRestoreTaskContext()`
- **Impact:** Complete blockage of FreeRTOS task execution

2 Experimental Methodology

2.1 System Configuration

- **Host Platform:** seL4 microkernel with CAmkES framework
- **Guest OS:** FreeRTOS 10.4.3 with ARM Cortex-A9 port

- **Architecture:** ARM32 (ARMv7-A) with virtualization extensions
- **Virtualization:** QEMU ARM virt machine, Cortex-A15 CPU
- **Memory Layout:** Base address 0x40000000, 512MB allocation

2.2 Integration Achievements

Prior to encountering the blocking issue, we successfully resolved multiple integration challenges:

1. **Architecture Consistency:** Resolved ARM32/AArch64 compilation mismatches
2. **Binary Format Compatibility:** Converted ELF to raw binary format for seL4 VM loader
3. **Memory Layout Alignment:** Configured consistent 0x40000000 base addressing
4. **Hardware Abstraction:**
 - UART communication (PL011 at 0x9000000)
 - GIC configuration (256-priority virtualized GIC at 0x8040000)
 - Memory management and heap allocation
5. **FreeRTOS Configuration:** Complete scheduler initialization and task creation

3 Detailed Technical Investigation

3.1 Context Switch Analysis

The failure occurs during the normal FreeRTOS scheduler startup sequence:

Listing 1: Normal Execution Flow

```

1 vTaskStartScheduler()
2   -> xPortStartScheduler()
3   -> vPortRestoreTaskContext()
4   -> portRESTORE_CONTEXT macro
5   -> RFEIA sp! instruction // FAILURE POINT

```

3.2 Assembly-Level Debugging

3.2.1 Stack Layout Verification

Our investigation confirmed that the task stack initialization is correct:

Listing 2: Stack Inspection Results

```

1 Stack pointer in TCB = 0x4000E0CC
2 stack[16] = 0x4000532C // Task function address (CORRECT)
3 stack[17] = 0x0000001F // System mode SPSR (CORRECT)

```

3.2.2 Memory Translation Testing

We verified that seL4's address translation is functioning correctly:

Listing 3: Address Translation Verification

```
1 === seL4 ADDRESS TRANSLATION CHECK ===
2 Current function address: 0x400075C0    // Valid range
3 Stack address range: TCB=0x4000E128    // Valid range
4 Test write/read to PC location: PASS    // Memory accessible
```

3.3 Root Cause Analysis

3.3.1 RFEIA Instruction Incompatibility

The RFEIA (Return From Exception In ARM state) instruction is designed for returning from interrupt/exception handlers. However, FreeRTOS uses this instruction for normal task startup, which creates a semantic mismatch in the virtualized environment:

- **Expected Context:** Exception return with saved processor state
- **Actual Context:** Fresh task startup with artificially created stack
- **seL4 Behavior:** Virtualization layer may restrict exception return semantics

3.3.2 Privilege Level Interaction

The consistent fault at PC: 0x8 (Software Interrupt vector) suggests that the RFEIA instruction may be triggering an unexpected exception or privilege violation within seL4's virtualization framework.

4 Experimental Evidence

4.1 Direct Function Call Test

To isolate the issue, we tested direct task function execution:

Listing 4: Direct Task Execution Test

```
1 // Direct function call bypassing context switch
2 extern void vPLCMain(void *pvParameters);
3 vPLCMain(NULL); // SUCCESSFUL EXECUTION
4 uart_puts("Direct call completed"); // This executes normally
```

Result: Task functions execute successfully when called directly, confirming that:

- Task code is correctly loaded and executable
- Memory layout and addressing work properly
- The issue is specifically in the context switch mechanism

4.2 Assembly Instruction Tracing

We implemented step-by-step debugging of the context restore process:

Listing 5: Custom Context Restore Debug (ARM Assembly)

```
1 vPortRestoreTaskContext:
2     CPS      #SYS_MODE           // Switch to system mode
3     LDR      R0, pxCurrentTCBConst
4     LDR      R1, [R0]
5     LDR      SP, [R1]             // Load task stack pointer
6     POP      {R1}                 // FPU context flag
7     POP      {R1}                 // Critical nesting
8     POP      {R0-R12, R14}        // Restore registers
9     LDR      R1, [SP]             // Load PC value
10    CMP      R1, #0x40000000       // Validate PC range
11    BLT      debug_invalid_pc     // Branch if invalid
12    BX       R1                   // Jump to task - FAULT OCCURS HERE
```

Observation: The system reaches the BX R1 instruction with a valid PC value (0x4000532C), but the execution results in a page fault at PC: 0x8.

5 Research Implications

5.1 Blocking Nature of the Issue

This problem represents a **complete blocker** for FreeRTOS execution under seL4:

- **No Task Execution:** FreeRTOS scheduler cannot start any tasks
- **No Workaround Available:** The context switch is fundamental to RTOS operation
- **Affects All FreeRTOS Applications:** Any application requiring task scheduling fails

5.2 Broader Impact on seL4 Virtualization

This issue may affect other operating systems that use similar ARM context switching mechanisms:

- Other real-time operating systems (e.g., Zephyr, ThreadX)
- Operating systems using ARM exception return instructions for task switching
- Any guest OS relying on specific ARM privilege mode semantics

6 Potential Research Directions

6.1 seL4 Virtualization Layer Investigation

1. **Exception Handling Analysis:** Investigate how seL4 handles ARM exception return instructions in virtualized contexts
2. **Privilege Mode Semantics:** Examine seL4's implementation of ARM privilege mode switches for guest OSes
3. **Hypervisor Trap Analysis:** Determine if specific ARM instructions are being trapped by seL4's virtualization layer

6.2 Alternative Context Switch Mechanisms

1. **Modified FreeRTOS Port:** Develop a seL4-specific FreeRTOS port that avoids problematic instructions
2. **Paravirtualization Approach:** Implement FreeRTOS context switching using seL4 system calls
3. **Cooperative Scheduling:** Investigate if cooperative (non-preemptive) scheduling works in seL4 environment

6.3 Formal Verification Considerations

1. **Verified Context Switch:** Research formally verified context switching mechanisms compatible with seL4
2. **Security Properties:** Analyze security implications of different context switch implementations
3. **Real-time Guarantees:** Ensure any solution maintains real-time scheduling properties

7 Technical Specifications

7.1 Error Context Details

Listing 6: Page Fault Context Dump

```
1 Pagefault from [vm0]: read prefetch fault @ PC: 0x8 IPA: 0x8, FSR: 0x6
2 Context:
3   pc: 0x8                                // ARM Software Interrupt vector
4   sp: 0x4000e0dc                         // Valid stack pointer
5   cpsr: 0x60000093                      // System mode, interrupts disabled
6   r1: 0x4000d15c                        // Contains task-related address
7   r14: 0x400021f8                       // Return address
```

7.2 Memory Layout Verification

- **Code Section:** 0x40000000 - 0x4001FFFF (128KB)
- **Data/BSS:** 0x40020000 - 0x4003FFFF (128KB)
- **Heap:** 0x40040000 - 0x4FFFFFFF (255MB)
- **Stack Region:** 0x4000C000 - 0x4000FFFF (16KB per task)

8 Conclusions

8.1 Critical Findings

1. The FreeRTOS-seL4 integration failure is **not due to configuration or setup issues**
2. The problem appears to be a **fundamental incompatibility** between FreeRTOS's ARM context switching and seL4's virtualization layer

3. **All other integration aspects work correctly**, including memory management, device I/O, and interrupt handling
4. The issue **completely blocks practical deployment** of FreeRTOS under seL4

8.2 Research Priority

This investigation reveals a critical research problem that requires deeper analysis of:

- seL4’s ARM virtualization implementation
- Interaction between guest OS context switching and hypervisor trap handling
- Alternative approaches for real-time OS virtualization on formally verified systems

8.3 Next Steps

1. **Collaborate with seL4 developers** to understand virtualization layer behavior
2. **Investigate seL4 source code** related to ARM exception handling and context switching
3. **Explore alternative RTOS solutions** that may be more compatible with seL4’s virtualization model
4. **Consider paravirtualized approaches** that work within seL4’s capability-based security model

9 References

- seL4 Reference Manual: <https://sel4.systems/Info/Docs/seL4-manual.pdf>
- FreeRTOS ARM Cortex-A Port Documentation
- ARM Architecture Reference Manual ARMv7-A
- CAmkES Component Architecture Framework Documentation